

Predictable Self-Organization with Computational Fields

Pt 2

Computational Fields: calculus, examples, self-stabilisation

Mirko Viroli¹ Jacob Beal²

¹ALMA MATER STUDIORUM—Università di Bologna, Italy
`mirko.viroli@unibo.it`

²BBN Technologies, USA
`jakebeal@bbn.com`

SASO 2014, London, 8/9/2014



- 1 Why a calculus of fields
- 2 The Field Calculus – “global viewpoint”
- 3 Semantic details – “local viewpoint”
- 4 Bootstrapping property verification: Self-stabilisation



- 1 Why a calculus of fields
- 2 The Field Calculus – “global viewpoint”
- 3 Semantic details – “local viewpoint”
- 4 Bootstrapping property verification: Self-stabilisation



Towards a foundational approach

- Many languages/modes incorporate features of aggregate programming
(see a survey in [Beal et al., 2013])
- No general formalisation approaches exist



Scope of the research

Towards a foundational approach

- Many languages/modes incorporate features of aggregate programming
(see a survey in [Beal et al., 2013])
- No general formalisation approaches exist

Our research steps

1. Identify a minimal set of ingredients of spatial computing
2. Devise a core calculus of computational fields [Viroli et al., 2013]
3. Provide a full formalisation in the proglangs/concurrency style
4. Isolate fragments enjoying certain properties




Perspective

- Paving the way towards advanced behavioural analysis
 - ▶ Sufficient conditions for self-stabilisation and density independence
 - ▶ Type-checking correctness properties
 - ▶ Formal local-to-global and global-to-local connections
- Devising tools for proper system engineering
 - ▶ building blocks, languages, patterns, APIs, infrastructure



An example: a “channel” deployed in physical space



- Computational field (or field): a mapping from nodes to values
- Can be: static, dynamic, or require time to stabilise
- Sometimes abstracted to a continuous space-time domain

How to achieve this structure in a self-organised way?

Local viewpoint

- What is the single-node behaviour?
- How should a node interact with neighbours and locally compute?

Global viewpoint

- What is the global structure that emerges, what are its properties?
- Can we create it compositionally?



How to achieve this structure in a self-organised way?

Local viewpoint

- What is the single-node behaviour?
- How should a node interact with neighbours and locally compute?

Global viewpoint

- What is the global structure that emerges, what are its properties?
- Can we create it compositionally?

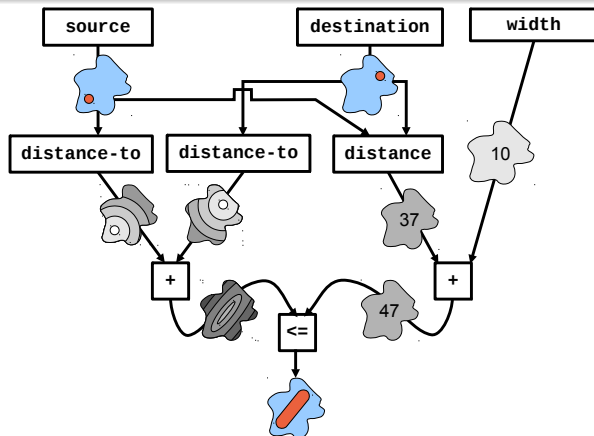
Computing with fields as first-class abstractions

- “Inputs” (dynamically coming from sensors) are fields!
 - “Outputs” (data produced, signal to actuators) are fields!
 - Computation is a function from inputs to outputs as usual
- ⇒ .. use a functional style to manipulate field expressions

How would we create a “channel” compositionally?

Using a functional style and mathematical operators

$\text{channel}(\text{src}, \text{dest}, \text{width}) =$
 $\text{dilate}(\text{distance-to}(\text{src}) + \text{distance-to}(\text{dest}) \leq \text{distance}(\text{src}, \text{dest}), \text{width})$



From global back to local viewpoint

Working at the global level is key

- Provides a more abstract framework
- Captures desired concepts more directly
- Facilitates design, specification, understanding

We need to map back to locality

- How do we “compile” the global specification back in the single node behaviour?

A core calculus of computational fields

What do we need to tackle the problem solidly?

- A handy formalisation framework
- Built on top of known and well-studied techniques
- Quickly leading to tool implementation



A core calculus of computational fields

What do we need to tackle the problem solidly?

- A handy formalisation framework
- Built on top of known and well-studied techniques
- Quickly leading to tool implementation

The notion of core calculus

- a tiny language: syntax + operational semantics (+ typing)
 - strives for the best tradeoff between:
 - ▶ compactness, simplicity (fewest constructs/concepts)
 - ▶ expressiveness (many program/behaviours)
 - popular examples:
 - ▶ λ -calculus [Barendregt, 1984]: a core for functional programming
 - ▶ π -calculus [Milner, 1999]: a core for interactive programming
 - ▶ FJ [Igarashi et al., 1999]: a core for object-oriented programming
- ⇒ virtually any programming constructs is formulated in that way..

The case of λ -calculus

Formalisation

- Syntax: lambdas, variables, application
- Semantics: one-step execution of function application

$$L ::= \lambda x.L \mid x \mid LL$$

$$(\lambda x.L)M \rightarrow L[M/x]$$



Outline

- 1 Why a calculus of fields
- 2 The Field Calculus – “global viewpoint”
- 3 Semantic details – “local viewpoint”
- 4 Bootstrapping property verification: Self-stabilisation



The calculus of computational fields [Viroli et al., 2013]

Main constructs

- Retains the functional style (recursive function definitions, one main)
- Built-in functions model sensor acquisition and any local algorithm
- Key computational field constructs
 - ▶ **rep** — Repetition: state evolving over time
 - ▶ **nbr** — Neighbouring: get information from neighbours
 - ▶ **if** — Restriction: a space-time branch

Other facts

- Syntax is LISP-like, in fact, a core of Proto [MIT Proto, 2012]
- Will use MIT Proto implementation for demoing
- Nodes all execute the same program, in asynchronous rounds
- Env. information (topology, sensors) extracted by built-in functions

$e ::=$	(field) expression:
l	local value (boolean, float, tuple ...)
$ x$	variable
$ (b \ e_1 \ e_2 \ \dots \ e_n)$	functional composition (b is built-in)
$ (f \ e_1 \ e_2 \ \dots \ e_n)$	function call (f is user-defined)
$ (\text{rep } x \ l \ e)$	time evolution
$ (\text{nbr } e)$	neighbourhood field construction
$ (\text{if } e_b \ e_t \ e_f)$	restriction

$F ::= (\text{def } f(x_1 \ x_2 \ \dots \ x_n) \ e)$ user-defined function




Functional composition: ($b \ e_1 \ \dots \ e_n$)

Goal: extending standard computation mechanisms to whole fields

- b is a built-in operator applied to $e_1 \ \dots \ e_n$ pointwise (node per node)
- a means to perform local operations: algorithms/sensing/acting
- they are pervasively used, shall use the following **color**

Basic example

- $(+ \ e_1 \ e_2)$: sums the two fields e_1 and e_2



Built-in functions as environment bindings

Examples

- (`red` 1) send value 1 to red actuator (a led)
- (`sense` "temperature") gives a field of temperature values
- (`sense` 1) gives a field of values from sensor n. 1
- (`dt`) gives the length of last round in each node
- (`mid`) gives the field of device identifiers



Built-in functions and tuples

How do we deal with more articulated data types

- they are still seen as locals
- built-in functions are ADT operations, they can be used to create and operate on such complex values

Examples using tuples

- (**tup** 10 20) gives everywhere a tuple of two values, 10 and 20
- (**1st** (**tup** 10 20)) accesses first component, i.e., 10
- (**2nd** (**tup** 10 20)) accesses second component, i.e., 20



Built-in functions and “special” neighbourhood fields

A peculiar data-type we handle is a neighbourhood field ϕ

A map $\{\sigma_1 \mapsto l_1, \dots, \sigma_n \mapsto l_n\}$ from neighbourhood to some local value

- its domain never escapes a node's neighbourhood
- it can't be the final result of computation, just an intermediate one
- some special built-in sensor could return one such field (`nbr-*`)
- some built-in function (`*-hood`, `*-hood+`) can turn it into a value

Examples using fields

- (`nbr-range`) gives a map ϕ from neighbours to their distance
- (`min-hood+ (nbr-range)`) gives the minimum of neighbour distances (+ means “myself excluded”)
- (`< (min-hood+ (nbr-range)) 5`) gives nodes with some short proximity

Function calls: ($f\ e_1 \dots e_n$)


Goal: providing standard name abstraction and recursive behaviour

- f is the name of a function to be defined as: (**def** $f(x_1 \dots x_n)\ e_{body}$)
- call-by-value semantics, possibly with recursion
- for user-defined functions shall use the following color

Example

Def (**def** **double**(x) ($*\ x\ 2$))

Use (**double** ($+\ e_1\ e_2$)): doubles the field ($+\ e_1\ e_2$)



$e ::=$

l
| x
| $(b\ e_1\ e_2\ \dots\ e_n)$
| $(f\ e_1\ e_2\ \dots\ e_n)$
| $(\text{rep}\ x\ l\ e)$
| $(\text{nbr}\ e)$
| $(\text{if}\ e_b\ e_t\ e_f)$

(field) expression:

local value (boolean, float, tuple ...)
variable
functional composition (b is built-in)
function call (f is user-defined)
time evolution
neighbourhood field construction
restriction

$F ::= (\text{def}\ f(x_1\ x_2\ \dots\ x_n)\ e)$

user-defined function




Field construct n.1: Time evolution (**rep** x 1 e)

Goal: supporting field evolution over time

- Initially it is globally 1 (or any expression..)
- At each new step a point in space updates to (local value of) e
- e can mention x, which stands for “previous field (state)”

Example

- (**rep** x 0 (+ x 1)): counts the number of rounds in each device
- (**rep** t 0 (+ t (dt))): the field of time




Field construct n.2: Neigh. field construction (**nbr** e)

Goal: enabling declarative node-to-node interaction

- The resulting field maps each node n to a neighbourhood field ϕ_n
- ϕ_n maps neighbours of n to their local value of e
- Is to be flattened by a \ast -hood built-in operator

Example

- Assume **bsns** be a special operator modelling a boolean sensor
- (**any-hood** (**nbr** (**bsns**))) : has any neighbour positive bsns?



Few field constructions examples

```
;; the field of number of neighbours
(def count-neighbours () (sum-hood (nbr 1)))

;; average distance of neighbours
(/ (sum-hood (nbr-range)) (count-neighbours))

;; most connected neighbour's id?
(def most-connected-neighbour ()
  (2nd (max-hood (nbr (tup (count-neighbours) (mid))))))
)

;; what's the highest connection in the network?
???
```



Neighbourhood chaining: nesting **nbr** into **rep**

```
;; what's the greatest value of F?  
(def goss-max (F) (rep x F (max-hood (nbr x))))  
  
;; which node is the highest connected?  
(def most-connected-node ()  
  (2nd (goss-max (tup (count-neighbours) (mid)))))  
  
;; what's the minimum hop-count distance to (any) source?  
(def hop-count (source)  
  (rep x (inf) (mux source 0 (min-hood+ (+ 1 (nbr x))))))  
  
;; distance-to, aka, gradient --> note it is robust and flexible!!  
(def distance-to (source)  
  (rep d (inf) (mux source 0 (min-hood+ (+ (nbr d) (nbr-range))))))
```



Building a channel

```
;; broadcast, aka gradcast, broadcasting the value of f at src
(def broadcast (src f)
  (rep t (tup (inf) f)
    (mux src (tup 0 f)
      (min-hood+ (tup (+ (nbr-range) (1st (nbr t)))
        (2nd (nbr t)))))))

;; minimum distance between a and b
(def distance (a b)
  (2nd (broadcast b (distance-to a))))

;; channel between src and dest of thickness width
(def channel (src dest width)
  (< (+ (distance-to src) (distance-to dest))
    (+ width (distance src dest))))
```




Field construct n.3: Restriction (**if** e_b e_t e_f)

Goal: providing a distributed branch

- e_b should be a boolean field (defining two restricted domains)
- Construct e_t where e_b is positive, and e_f where e_b is negative
- Not a mere superposition of e_t and e_f , but a true domain restriction!
 - ▶ i.e., nbrs inside e_t should not escape into where e_b is false

Example

- $(\text{if } (\text{bsns}) \text{ } e \text{ } 0)$ creates e where bsns is true, and 0 elsewhere




Channel avoiding obstacles

```
;; automatically circumventing the obstacle
(def channel-avoiding-obstacle (obstacle source destination width)
  (if (not obstacle) (channel source destination width) 0))

;; mux applies the obstacle
;; but it does not interfere with channel computation
(def wrong-channel-avoiding-obstacles (obstacle source destination width)
  (mux (not obstacle) (channel source destination width) 0))
```



Channel in action – a simulation in Proto



Outline

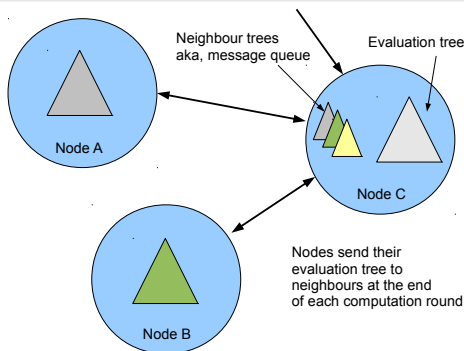
- 1 Why a calculus of fields
- 2 The Field Calculus – “global viewpoint”
- 3 Semantic details – “local viewpoint”
- 4 Bootstrapping property verification: Self-stabilisation



Key aspects of the semantics: network model

Network model

- A node has a state S (evaluation tree) updated at asynchronous rounds
- At the end of the round, S is spread to the (current) neighbourhood
- State is updated “against” the neighbour states just received



Key aspects of the semantics: node (abstract) model

Inside a node: the evaluation tree S

- S is an annotated version of the expression to evaluate
 - ▶ annotations are used to keep track of the next evaluation site
 - ▶ some annotations are persistent, and used to interact in space/time
- S may dynamically expand due to (recursive) calls
- field constructs semantics impact shape of annotations
 - ▶ `rep`: last result is stored in an annotation and reminded at next round
 - ▶ `nbr`: observes annotations in the same position in neighbour trees
 - ▶ `if`: discards the neighbour trees which took a different branch

Main issue (and contribution wrt previous formalisation attempts)

Accommodate the interplay between `rep`, `if` and `nbr` – even in the presence of (recursive) calls



A node state is an annotated evaluation tree

Round 1 $(\text{rep } x \ 0 \ (+ \ \underline{x} \ 1)) \rightarrow (\text{rep } x \ 0 \ (+ \ x \cdot \underline{0} \ \underline{1})) \rightarrow$
 $(\text{rep } x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot \underline{1})) \rightarrow \underline{(\text{rep } x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot 1) \cdot \underline{1})}$
 $\rightarrow (\text{rep}^1 x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot 1) \cdot 1) \cdot \underline{1}$



A node state is an annotated evaluation tree

Round 1 $(\text{rep } x \ 0 \ (+ \ \underline{x} \ 1)) \rightarrow (\text{rep } x \ 0 \ (+ \ x \cdot 0 \ \underline{1})) \rightarrow$
 $(\text{rep } x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot 1)) \rightarrow \underline{(\text{rep } x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot 1) \cdot 1)}$
 $\rightarrow (\text{rep}^1 x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot 1) \cdot 1) \cdot 1$

Round 2 $(\text{rep}^1 x \ 0 \ (+ \ x \ 1)) \rightarrow .. \rightarrow (\text{rep}^2 x \ 0 \ (+ \ x \cdot 1 \ 1 \cdot 1) \cdot 2) \cdot 2$



A node state is an annotated evaluation tree

Round 1 $(\text{rep } x \ 0 \ (+ \ \underline{x} \ 1)) \rightarrow (\text{rep } x \ 0 \ (+ \ x \cdot \underline{0} \ \underline{1})) \rightarrow$
 $(\text{rep } x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot \underline{1})) \rightarrow \underline{(\text{rep } x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot 1) \cdot \underline{1})}$
 $\rightarrow (\text{rep}^1 x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot 1) \cdot 1) \cdot \underline{1}$

Round 2 $(\text{rep}^1 x \ 0 \ (+ \ x \ 1)) \rightarrow .. \rightarrow (\text{rep}^2 x \ 0 \ (+ \ x \cdot 1 \ 1 \cdot 1) \cdot 2) \cdot \underline{2}$

Round 3 $(\text{rep}^2 x \ 0 \ (+ \ x \ 1)) \rightarrow .. \rightarrow (\text{rep}^3 x \ 0 \ (+ \ x \cdot 2 \ 1 \cdot 1) \cdot 3) \cdot \underline{3}$




A node state is an annotated evaluation tree

Round 1 $(\text{rep } x \ 0 \ (+ \ \underline{x} \ 1)) \rightarrow (\text{rep } x \ 0 \ (+ \ x \cdot 0 \ \underline{1})) \rightarrow$
 $(\text{rep } x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot \underline{1})) \rightarrow (\text{rep } x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot 1) \cdot \underline{1})$
 $\rightarrow (\text{rep}^1 x \ 0 \ (+ \ x \cdot 0 \ 1 \cdot 1) \cdot 1)$

Round 2 $(\text{rep}^1 x \ 0 \ (+ \ x \ 1)) \rightarrow .. \rightarrow (\text{rep}^2 x \ 0 \ (+ \ x \cdot 1 \ 1 \cdot 1) \cdot 2) \cdot \underline{2}$

Round 3 $(\text{rep}^2 x \ 0 \ (+ \ x \ 1)) \rightarrow .. \rightarrow (\text{rep}^3 x \ 0 \ (+ \ x \cdot 2 \ 1 \cdot 1) \cdot 3) \cdot \underline{3}$



Evaluation aligns with neighbours to support nbr

Neighbour trees:

$$\sigma_1 \mapsto (\text{min-hood } (\text{nbr } (\text{sns}) \cdot 4) \cdot \phi_1) \cdot 4$$

$$\sigma_2 \mapsto (\text{min-hood } (\text{nbr } (\text{sns}) \cdot 9) \cdot \phi_2) \cdot 9$$

Evaluation (assume here sns gives 7):

$$(\text{min-hood } (\text{nbr } (\text{sns}))) \rightarrow$$

$$(\text{min-hood } (\text{nbr } (\text{sns}) \cdot 7)) \rightarrow$$

$$(\text{min-hood } (\text{nbr } (\text{sns}) \cdot 7) \cdot (\sigma \mapsto 7, \sigma_1 \mapsto 4, \sigma_2 \mapsto 9)) \rightarrow$$

$$\frac{(\text{min-hood } (\text{nbr } (\text{sns}) \cdot 7) \cdot (\sigma \mapsto 7, \sigma_1 \mapsto 4, \sigma_2 \mapsto 9))}{(\text{min-hood } (\text{nbr } (\text{sns}) \cdot 7) \cdot (\sigma \mapsto 7, \sigma_1 \mapsto 4, \sigma_2 \mapsto 9)) \cdot 4}$$



Evaluation aligns with neighbours to support nbr

Neighbour trees:

$$\sigma_1 \mapsto (\text{min-hood } (\text{nbr } (\text{sns}) \cdot 4) \cdot \phi_1) \cdot 4$$

$$\sigma_2 \mapsto (\text{min-hood } (\text{nbr } (\text{sns}) \cdot 9) \cdot \phi_2) \cdot 9$$


Evaluation (assume here sns gives 7):

$$(\text{min-hood } (\text{nbr } (\text{sns}))) \rightarrow$$

$$(\text{min-hood } (\text{nbr } (\text{sns}) \cdot 7)) \rightarrow$$

$$(\text{min-hood } (\text{nbr } (\text{sns}) \cdot 7) \cdot (\sigma \mapsto 7, \sigma_1 \mapsto 4, \sigma_2 \mapsto 9)) \rightarrow$$

$$(\text{min-hood } (\text{nbr } (\text{sns}) \cdot 7) \cdot (\sigma \mapsto 7, \sigma_1 \mapsto 4, \sigma_2 \mapsto 9)) \cdot 4$$



Restriction prevents escaping a domain

Neighbour trees:

$$\sigma_1 \mapsto (\text{if } (\text{bsns}) \cdot f \text{ (min-hood (nbr (sns))) } 0 \cdot 0) \cdot 0$$

$$\sigma_2 \mapsto (\text{if } (\text{bsns}) \cdot t \text{ (min-hood (nbr (sns) \cdot 9) \cdot \phi_2 \ 0)) \cdot 9$$

Evaluation (assume here sns gives 7 and bsns gives t):

$$(\text{if } \underline{(\text{bsns})} \text{ (min-hood (nbr (sns))) } 0) \rightarrow$$

$$(\text{if } (\text{bsns}) \cdot t \text{ (min-hood (nbr } \underline{(\text{sns})})} 0) \rightarrow \dots \rightarrow$$

$$\underline{(\text{if } (\text{bsns}) \cdot t \text{ (min-hood (nbr (sns) \cdot 7) \cdot (\sigma \mapsto 7, \sigma_2 \mapsto 9)) \cdot 7 \ 0)}$$



Restriction prevents escaping a domain

Neighbour trees:

$$\sigma_1 \mapsto (\text{if } (\text{bsns}) \cdot f \text{ (min-hood (nbr (sns))) } 0 \cdot 0) \cdot 0$$


$$\sigma_2 \mapsto (\text{if } (\text{bsns}) \cdot t \text{ (min-hood (nbr (sns) \cdot 9) \cdot \phi_2 \ 0)) \cdot 9$$

Evaluation (assume here sns gives 7 and bsns gives t):

$$(\text{if } (\text{bsns}) \text{ (min-hood (nbr (sns))) } 0) \rightarrow$$

$$(\text{if } (\text{bsns}) \cdot t \text{ (min-hood (nbr (sns))) } 0) \rightarrow \dots \rightarrow$$

$$(\text{if } (\text{bsns}) \cdot t \text{ (min-hood (nbr (sns) \cdot 7) \cdot (\sigma \mapsto 7, \sigma_2 \mapsto 9)) \cdot 7 \ 0)$$



Recursively, neighbour trees that are not aligned are temporarily discarded



Function calls dynamically expand the evaluation tree

Take definition (**def** fun (x) (min-hood (nbr x)))..


When evaluating (if (bsns) (fun sns) 0), function body is expanded:

$$\begin{aligned} &(\text{if } \underline{\text{bsns}} \text{ (fun sns) 0}) \rightarrow (\text{if } (\text{bsns}) \cdot t \text{ (fun } \underline{\text{sns}}) \text{ 0}) \rightarrow \\ &(\text{if } (\text{bsns}) \cdot t \text{ } \underline{\text{(fun (sns) } \cdot 7 \text{) 0}}) \rightarrow \\ &(\text{if } (\text{bsns}) \cdot t \text{ (fun}^{(\text{min-hood (nbr } \underline{x})})} \text{ (sns) } \cdot 7 \text{) 0}) \rightarrow \\ &(\text{if } (\text{bsns}) \cdot t \text{ (fun}^{(\text{min-hood (nbr } \underline{x} \cdot 7 \text{)})} \text{ (sns) } \cdot 7 \text{) 0}) \rightarrow \dots \end{aligned}$$


Function calls dynamically expand the evaluation tree

Take definition (**def** fun (x) (min-hood (nbr x)))..

When evaluating (if (bsns) (fun sns) 0), function body is expanded:

$$\begin{aligned} &(\text{if } \underline{(\text{bsns})} \text{ (fun sns) } 0) \rightarrow (\text{if } (\text{bsns}) \cdot t \text{ (fun } \underline{(\text{sns})}) \text{ } 0) \rightarrow \\ &(\text{if } (\text{bsns}) \cdot t \text{ (fun } \underline{(\text{sns})} \cdot 7) \text{ } 0) \rightarrow \\ &(\text{if } (\text{bsns}) \cdot t \text{ (fun } ^{(\text{min-hood (nbr x)})} \text{ (sns)} \cdot 7) \text{ } 0) \rightarrow \\ &(\text{if } (\text{bsns}) \cdot t \text{ (fun } ^{(\text{min-hood (nbr x} \cdot 7))} \text{ (sns)} \cdot 7) \text{ } 0) \rightarrow \dots \end{aligned}$$


Expansion at function call, contraction on non-taken branches

The pillars of the operational semantics

Elements


- A node's computation is about annotating the evaluation tree
- Such a tree is spread to neighbours after being cleaned up
- Neighbour trees affect evaluation
- Annotations to **nbr**
 - ▶ persist over “space” (sent to neighbours)
 - ▶ used to remotely reconstruct a neighbourhood field at **nbr** sites
- Annotations to **if**
 - ▶ persist over “space” (sent to neighbours)
 - ▶ used to filter out neighbours that do not match at **if** sites
- Annotations to **rep**
 - ▶ persist over “time” (remembered at next round)
 - ▶ used to recompute at **rep** sites



Tree evaluation: pictorial semantics

Neighbour trees


Evaluation tree



Tree evaluation: pictorial semantics

Neighbour trees

Evaluation tree



The operational semantics of node model

Runtime Expression Syntax:

$e ::= a\bar{v}$	runtime expression (rte)
$a ::= x \mid v \mid (\text{nbr } e) \mid (\text{if } eee) \mid (\text{rep}^s x w e) \mid (f^s \bar{e}) \mid (o \bar{e})$	auxiliary rte
$v ::= 1 \mid \phi$	runtime value
$s ::= \hat{a}$	superscript
$w ::= x \mid 1$	variable or local value
$\phi ::= \bar{\sigma} \mapsto \bar{1}$	field value
$\Theta ::= \bar{\sigma} \mapsto \bar{e}$	tree environment
$\Gamma ::= \bar{x} := \bar{v}$	variable environment

Congruence Contexts:

$$\mathbb{C} ::= (\text{nbr } \square) \mid (f^s \bar{e} \square \bar{e}) \mid (o \bar{e} \square \bar{e}) \mid (\text{if } \square e e) \mid (\text{if } at \square e) \mid (\text{if } af e \square)$$

Alignment contexts:

$$\mathbb{A} ::= \mathbb{C} \mid (\text{rep}^s x w \square) \mid (f \square \bar{a}\bar{v})$$

Auxiliary functions:

$$\begin{aligned} \pi_{\mathbb{A}}(\Theta, \Theta') &= \pi_{\mathbb{A}}(\Theta), \pi_{\mathbb{A}}(\Theta') \\ \pi_{\mathbb{A}}(\sigma \mapsto (\mathbb{A}'[e])v) &= \sigma \mapsto e \quad \text{if } \mathbb{A}' :: \mathbb{A} \\ \pi_{\mathbb{A}}(\sigma \mapsto e) &= \bullet \quad \text{otherwise} \end{aligned}$$

$$\begin{aligned} s \triangleright a &= a \\ s \triangleright o &= s \end{aligned}$$

$$\begin{aligned} (\text{nbr } \square) &:: (\text{nbr } \square) \\ (f^{s'} e'_1 \dots e'_{l-1} \square e'_{l+1} \dots e'_n) &:: (f^s e_1 \dots e_{l-1} \square e_{l+1} \dots e_n) \\ (o e'_1 \dots e'_{l-1} \square e'_{l+1} \dots e'_n) &:: (o e_1 \dots e_{l-1} \square e_{l+1} \dots e_n) \\ (\text{if } \square e'_1 e'_2) &:: (\text{if } \square e_1 e_2) \\ (\text{if } at \square e') &:: (\text{if } at \square e) \\ (\text{if } af e' \square) &:: (\text{if } af e \square) \\ (\text{rep}^s x w \square) &:: (\text{rep}^s x w \square) \\ (f \square e'_1 \dots e'_n) &:: (f \square e_1 \dots e_n) \end{aligned}$$

Reduction Rules:

$$\begin{aligned} &\frac{[\text{THEN}]}{\Theta; \Gamma \vdash (\text{if } at \ a \ 1 \ e) \rightarrow (\text{if } at \ a \ 1 \ |e|)1} \\ &\frac{[\text{VAL}]}{\Theta; \Gamma \vdash v \rightarrow v^v} \quad \frac{[\text{ELSE}]}{\Theta; \Gamma \vdash (\text{if } af \ e \ a \ 1) \rightarrow (\text{if } af \ |e| \ a \ 1)1} \\ &\frac{[\text{VAR}]}{\Theta; \Gamma \vdash x \rightarrow x\Gamma(x)|_{\text{dom}(\Theta), \mathcal{E}(\text{self})}} \quad \frac{[\text{CONG}]}{\Theta; \Gamma \vdash \mathbb{C}[a] \rightarrow \mathbb{C}[e]} \\ &\frac{[\text{NBR}]}{\Theta; \Gamma \vdash (\text{nbr } a \ 1) \rightarrow (\text{nbr } a \ 1)\phi} \quad \frac{[\text{REP}]}{\Theta; \Gamma \vdash (\text{rep}^1 x w \square) (\Theta); \Gamma, (x := (\Gamma(w) \triangleright \hat{1})) \vdash a \rightarrow a^{\hat{v}}} \\ &\frac{[\text{OP}]}{\Theta; \Gamma \vdash (o \ \bar{a}\bar{v}) \rightarrow (o \ \bar{a}\bar{v})\mathcal{E}(o, \bar{v})} \quad \frac{[\text{FUN}]}{\Theta; \Gamma \vdash (f^a \ \bar{a}\bar{v}) \rightarrow (f^a \ \bar{a}\bar{v})^{\hat{v}}} \end{aligned}$$



Outline

- 1 Why a calculus of fields
- 2 The Field Calculus – “global viewpoint”
- 3 Semantic details – “local viewpoint”
- 4 Bootstrapping property verification: Self-stabilisation



Bootstrapping property verification

What is the field calculus good at?

- Giving feedbacks on existing/new models/implementations
 - Helping designing new features (e.g., typing, higher-order functions)
 - Assisting development of surface languages, APIs
- ⇒ Predicting the behaviour of fragments of the language



Bootstrapping property verification

What is the field calculus good at?

- Giving feedbacks on existing/new models/implementations
 - Helping designing new features (e.g., typing, higher-order functions)
 - Assisting development of surface languages, APIs
- ⇒ Predicting the behaviour of fragments of the language

History and Roadmap

- PAST: Started in SAPERE EU FP7 project [Zambonelli et al., 2011]
 - ▶ self-stabilisation for a rule-based field language
- PRESENT: A self-stabilisation result for a fragment of the field calculus achieved [Viroli and Damiani, 2014]
- PRESENT: Universality of field calculus [Beal et al., 2014]
- ONGOING: Density-independence / approximability
- FUTURE: Extending the fragments of investigation

Self-stabilisation

Traditional self-stabilisation [Dolev, 2000]

A distributed system that is self-stabilising will end up in a *correct* state (in finite steps) no matter what state it is initialised with.

Superstabilisation

A distributed system is superstabilising if it is self-stabilising and it recovers *fast* from (single) topological changes.

Considerations for the field-calculus

- Self-stabilisation is of course a key fault-tolerant property
- The definition should somehow be adapted to work with fields
- One still needs to (implicitly) define what are the correct states
- Knowing that (and which) correct states are reached is indeed about predictability!
- What about: (**def** **flip** (a b) (**rep** x a (- (+ a b) x)))

Self-stabilisation for fields [Viroli and Damiani, 2014]

We considered a notion of (*unique,super-*) self-stabilisation

If a field expression is self-stabilising, then it necessarily reaches a **unique** configuration (unique fixpoint), recovering from *any* change

- assume execution of rounds is asynchronous but *fair*
- fixpoint is independent of initial state
- fixpoint is reached in finite time
- fixpoint depends on environment E (sensors/topology)
- if E (arbitrarily) changes the system recovers

Implication

Any self-stabilising field expression e has denotational semantics:

$$\phi_{final} = \phi_E(e)$$

..i.e., a program e has predictable global/final outcome ϕ_{final}

Sufficient conditions: a self-stabilising fragment

Ideas

- Generalise over a key block, the hop-count distance (**gradient**):

```
(rep x (inf) (mux source 0 (min-hood+ (+ 1 (nbr x))))
```

- **nbr** should stay inside a ***-hood**, itself inside a **rep**
- ***-hood** should implement an order-independent function (e.g. **min**)
- Should generalise over the “+ 1” function
- Should guarantee that propagation **terminates**
- Strict layering of self-stabilising fields is self-stabilising



Some preliminaries

Rework a distance-to (gradient) specification

⇒ `(rep x (inf) (mux source 0 (min-hood+ (+ 1 (nbr x))))))`

- .. assume source is 0 on sources, and ∞ elsewhere

⇒ `(rep x (inf) (min source (min-hood+ (+ 1 (nbr x))))))`

- .. generalise over function `+`, which could also use sensor-dependent expressions

⇒ `(def spr-aggr (E0 FUN E1 .. En)
 (rep x (inf)
 (min E0 (FUN (min-hood+ x) E1 .. En))))`

- .. abstract this block into a syntactic construct

⇒ $\{e_0 : g(@, e_1, \dots, e_n)\}$



The Syntax – 3 main ingredients

$e ::= x$		variable
v		value
s		<u>sensor</u>
$g(e_1, \dots, e_n)$		<u>pointwise functional composition</u>
$\{e : g(@, e_1, \dots, e_n)\}$		<u>spreading-aggregation</u> (nbr,rep,min-hood+)
$g ::= f$		user-defined function
b		built-in function

Additional constraints

- User-defined function definitions have no cycles
- Built-in functions are environment independent
- Basic typing properties are satisfied



Stabilising progressions

A function g over type T is a *stabilising progression* if:


- (i) g is a *pure operator* — it calls no sensor or spreading-aggregation;
- (ii) T is *locally noetherian* — $\llbracket T \rrbracket$ is equipped with a total order relation \leq_T , and for every element $v \in \llbracket T \rrbracket$, there are no infinite ascending chains of elements $v_0 <_T v_1 <_T v_2 \cdots$ such that (for every $n \geq 0$) $v_n <_T v$;
- (iii) g is *monotone* in its first argument — $v \leq_T v'$ implies $\llbracket g \rrbracket(v, \bar{v}) \leq_T \llbracket g \rrbracket(v', \bar{v})$ for any \bar{v} ;
- (iv) g is *progressive* in its first argument — $v <_T \llbracket g \rrbracket(v, \bar{v})$ (for $v \neq \text{top}(T)$).

Self-stabilisation result

A field expression whose spreading aggregation constructs feature only self-stabilising progressions is self-stabilising


Distance-to

```
def int distanceto(int i) is { i : @ + #dist }
```




Distance-to cirumventing an obstacle

```
def int distanceto(int i) is { i : @ + #dist }  
def int restrict(int i, bool b) is b ? i : INF  
def int distobs(int i, bool b) is { i : restrict(@ + #dist, b) }
```



Channel

```
def int distanceto(int i) is { i : @ + #dist }
def int restrict(int i, bool b) is b ? i : INF
def int distobs(int i, bool b) is { i : restrict(@ + #dist, b) }
def <int,int> add_to_1st(<int,int> x, int y) is <1st(x)+ y, 2nd(x)>
def <int,int> broadcast(int i, int j) is { <i, j>: add_to_1st(@, #dist) }
def int dist(int i, int j) is broadcast(restrict(j,j==0),distanceto(i))
def bool path(int i, int j, int w) is
    distanceto(i) + distanceto(j) < dist(i,j) + w
def int channel(int i, int j, int w) is
    broadcast(distanceto(j),not path(i,j,w))
```



The proof

Structure of the proof

- Because of layering can reason inductively on expression structure
- The only non-trivial case is to prove that $\{v_0 : g(@, v_1, \dots, v_n)\}$ self-stabilises (v_i are immutable fields)

Key idea

Inductively on the size of the subnetwork S that already self-stabilised

- Base case: the node holding the minimum value of v_0 self-stabilises immediately (at its first round)
- Inductive case: after a small transient, the minimum value outside S necessarily increases until $S \neq \emptyset$, so there's surely another node that will self-stabilise:
 - ▶ either one reaching $\text{top}(T)$,
 - ▶ or one in the neighbourhood of a node in S... in both cases the result is independent of the initial state

Non-self-stabilising fields

```
;; previous definition
(def goss-max (F) (rep x F (max-hood (nbr x))))

;; non-self-stabilising
(goss-max (sense 1))

;; still not working
(def goss-max2 (F) (rep x F (max-hood (nbr x))))

;; where's the bug?
(def wrong-distance-to (SRC)
  (rep x (inf) (mux SRC 0 (min-hood (+ (nbr-range) (nbr x))))))
```



Conclusions

Other predictability studies (in the following)

- Eventual consistency, i.e., network density independence
- Universality of field calculus constructs

Open issues

- There are more self-stabilising fields than our condition can check
- The notion of self-stabilisation can be generalised
- What about correct dynamic fields?
- There's room for some more foundational work
- You do not like the language itself? Others are under study..



References I

- [Barendregt, 1984] Barendregt, H. P. (1984).
The Lambda Calculus.
North Holland, revised edition.
- [Beal et al., 2013] Beal, J., Dulman, S., Usbeck, K., Viroli, M., and Correll, N. (2013).
Organizing the aggregate: Languages for spatial computing.
In Mernik, M., editor, *Formal and Practical Aspects of Domain-Specific Languages: Recent Developments*, chapter 16, pages 436–501. IGI Global.
A longer version available at: <http://arxiv.org/abs/1202.5509>.
- [Beal et al., 2014] Beal, J., Viroli, M., and Damiani, F. (2014).
Towards a unified model of spatial computing.
In *7th Spatial Computing Workshop (SCW 2014)*, AAMAS 2014, Paris, France.
- [Dolev, 2000] Dolev, S. (2000).
Self-Stabilization.
MIT Press.
- [Igarashi et al., 1999] Igarashi, A., Pierce, B., and Wadler, P. (1999).
Featherweight Java: A minimal core calculus for Java and GJ.
- [Milner, 1999] Milner, R. (1999).
Communicating and Mobile Systems: The π -calculus.
Cambridge University Press.
- [MIT Proto, 2012] MIT Proto (Retrieved June 1st, 2012).
MIT Proto.
software available at <http://proto.bbn.com/>.



References II

[Viroli and Damiani, 2014] Viroli, M. and Damiani, F. (2014).

A calculus of self-stabilising computational fields.

In Eva Kühn and Pugliese, R., editors, *Coordination Languages and Models*, volume 8459 of *LNCS*, pages 163–178.

Springer-Verlag.

Proceedings of the 16th Conference on Coordination Models and Languages (Coordination 2014), Berlin (Germany), 3-5 June.

[Viroli et al., 2013] Viroli, M., Damiani, F., and Beal, J. (2013).

A calculus of computational fields.

In Canal, C. and Villari, M., editors, *Advances in Service-Oriented and Cloud Computing*, volume 393 of *Communications in Computer and Information Science*, pages 114–128. Springer Berlin Heidelberg.

[Zambonelli et al., 2011] Zambonelli, F., Castelli, G., Ferrari, L., Mamei, M., Rosi, A., Serugendo, G. D. M., Risoldi, M., Tchao, A.-E., Dobson, S., Stevenson, G., Ye, J., Nardini, E., Omicini, A., Montagna, S., Viroli, M., Ferscha, A., Maschek, S., and Wally, B. (2011).

Self-aware pervasive service ecosystems.

Procedia CS, 7.

